

Generative Adversarial Networks

Moez Krichen

ReDCAD, University of Sfax, Sfax, Tunisia
moez.krichen@redcad.org

Abstract—Generative Adversarial Networks (GANs) are a type of deep learning techniques that have shown remarkable success in generating realistic images, videos, and other types of data. This paper provides a comprehensive guide to GANs, covering their architecture, loss functions, training methods, applications, evaluation metrics, challenges, and future directions. We begin with an introduction to GANs and their historical development, followed by a review of the background and related work. We then provide a detailed overview of the GAN architecture, including the generator and discriminator networks, and discuss the key design choices and variations. Next, we review the loss functions utilized in GANs, including the original minimax objective, as well as more recent approaches s.a. Wasserstein distance and gradient penalty. We then delve into the training of GANs, discussing common techniques s.a. alternating optimization, minibatch discrimination, and spectral normalization. We also provide a survey of the various applications of GANs across domains. In addition, we review the evaluation metrics utilized to assess the diversity and quality of GAN-produced data. Furthermore, we discuss the challenges and open issues in GANs, including mode collapse, training instability, and ethical considerations. Finally, we provide a glimpse into the future directions of GAN research, including improving scalability, developing new architectures, incorporating domain knowledge, and exploring new applications. Overall, this paper serves as a comprehensive guide to GANs, providing both theoretical and practical insights for researchers and practitioners in the field.

Index Terms—Generative Adversarial Networks, GANs, Applications, Evaluation metrics, Limitations.

I. INTRODUCTION

The rapid development of the disciplines of machine learning (ML) and deep learning (DL) has led to significant advances in artificial intelligence (AI) [1]. The goal of ML, a branch of artificial intelligence, is to create algorithms that can automatically learn from data and generate predictions or judgments. Figure 1 shows how traditional programming and machine learning vary from each other.

In contrast, DL is a subset of ML that utilizes neural networks (NNs) with multiple layers to learn complex data representations [2]. As a form of DL, GANs have demonstrated remarkable success in producing high-quality data and have become a crucial instrument for data synthesis and enhancement.

Due to its capacity to learn complex data representations, DL has revolutionized a variety of disciplines, including computer vision, NLP¹, and speech recognition. DL models, including CNNs² and RNNs³, have attained state-of-the-art




Fig. 1. Classical Programming vs ML.

performance on a variety of difficult tasks, including image recognition and language translation. GANs, which employ a specific form of NN architecture, have created new opportunities for data generation and synthesis and have been utilized to produce realistic images, videos, and audio.

In this paper, we focus specifically on Generative Adversarial Networks (GANs), which are a form of DL, and provide a comprehensive guide to their architecture, loss functions, training methods, applications, evaluation metrics, challenges, and future directions [3]. We believe that GANs represent a significant advancement in the field of AI, and have the potential to unlock new opportunities for scientific discovery and artistic expression. By providing a thorough overview of GANs, we hope to make this complex and challenging topic more accessible to researchers and practitioners in the field, and inspire new innovations and applications.

GANs have emerged as a useful technique for producing realistic data in a variety of disciplines ranging from computer vision and graphics to NLP and audio synthesis. Goodfellow et al. developed GANs in 2014, and they have since been one of the most active areas of research in DL. GANs are made up of two adversarial-trained NNs, a generator network ($Generator_{Network}$) and a discriminator network ($Discriminator_{Network}$). The $Generator_{Network}$ learns to produce data that is indistinguishable from the real data ($Real_{Data}$), whereas the $Discriminator_{Network}$ learns to differentiate between the two.

Despite their success, GANs are a complex and challenging topic that require a deep understanding of both the underlying theory and practical implementation. In this paper, we provide a comprehensive guide to GANs, covering their architecture, loss functions, training methods, applications, evaluation metrics, challenges, and future directions. Our goal is to provide

¹NLP=Natural Language Processing

²CNN=Convolutional Neural Network

³RNN=Recurrent Neural Network

both theoretical and practical insights for researchers and practitioners in the field, and to help demystify the often confusing and intimidating aspects of GANs.

The primary contributions of this paper are:

- 1) A comprehensive overview of the GAN architecture, including the *Generator_{Network}* and *Discriminator_{Network}*, and the key design choices and variations.
- 2) An in-depth review of the loss functions utilized in GANs, including the original minimax objective, as well as more recent approaches s.a. Wasserstein distance and gradient penalty.
- 3) A survey of the various training methods utilized in GANs, including alternating optimization, minibatch discrimination, and spectral normalization.
- 4) A review of the different applications of GANs across domains s.a. computer vision, NLP, and audio synthesis.
- 5) An exploration of the evaluation metrics utilized to assess the diversity and quality of GAN-produced data.
- 6) A discussion of the challenges and open issues in GAN research, including training instability, mode collapse, and ethical considerations.
- 7) A glimpse into the future directions of GAN research, including improving scalability, developing new architectures, incorporating domain knowledge, and exploring new applications.

In Section II, we provide a brief background on GANs and related work. In Section III, we provide a detailed overview of the GAN architecture, including the *Generator_{Network}* and *Discriminator_{Network}*, and the key design choices and variations. In Section IV, we review the loss functions utilized in GANs, including the original minimax objective, as well as more recent approaches s.a. Wasserstein distance and gradient penalty. In Section V, we discuss the training methods utilized in GANs, including alternating optimization, minibatch discrimination, and spectral normalization. In Section VI, we survey the various applications of GANs across domains s.a. computer vision, NLP, and audio synthesis. In Section VII, we explore the evaluation metrics utilized to assess the diversity and quality of GAN-produced data, s.a. Frechet Inception Distance and Inception Score. In Section VIII, we discuss the challenges and open issues in GAN research, including training instability, mode collapse, and ethical considerations. Finally, in Section IX, we provide a glimpse into the future directions of GAN research, including improving scalability, developing new architectures, incorporating domain knowledge, and exploring new applications. We conclude the paper in Section X, summarizing our contributions and discussing the broader impact and potential of GANs.

II. BACKGROUND

In recent years, DL has emerged as an important tool for solving a wide range of ML problems, s.a. image classification, speech recognition, and NLP [4], [5]. DL algorithms are based on NNs, which are composed of layers of interconnected

processing nodes that can learn to recognize patterns in data through a process of supervised or unsupervised training.

Generative models are a class of DL algorithm that can produce new data that is similar to the *TrainingData* (*TrainingData*). They have many types of applications, from image synthesis to speech generation. One of the most popular types of generative models is GAN.

The basic concept of GANs was introduced by Ian Goodfellow and his colleagues in 2014 [6]. As illustrated in Figure 2 GANs consist of two NNs: a *Generator_{Network}* and a *Discriminator_{Network}*. The *Generator_{Network}* takes as input a random noise vector and produces a new sample that is intended to be similar to the *TrainingData*. The *Discriminator_{Network}* takes as input a sample and tries to differentiate between samples produced by the *Generator_{Network}* and samples from the *TrainingData*. The *Generator_{Network}* is trained to produce samples that are difficult for the *Discriminator_{Network}* to differentiate from the *TrainingData*, while the *Discriminator_{Network}* is trained to classify samples correctly as either real or fake. The training process (*TrainingProcess*) for GANs is iterative, and involves alternating between training the *Generator_{Network}* and *Discriminator_{Network}*. During training, the *Generator_{Network}* learns to produce more realistic samples, while the *Discriminator_{Network}* learns to become more accurate in differentiating between real and fake samples. The goal is to find an equilibrium where the *Generator_{Network}* produces samples that are indistinguishable from the *TrainingData*, and the *Discriminator_{Network}* is not able to differentiate between real and fake samples.

Several types of GAN architectures have been proposed, including deep convolutional GANs, Wasserstein GANs [7], and conditional GANs [8]. DCGANs are a type of GAN that use CNNs in the *Generator_{Network}* and *Discriminator_{Network}* to produce high-quality images. WGANs are a type of GAN that use the Wasserstein distance metric instead of the traditional Jensen-Shannon divergence to evaluate the distance between the produced and real distributions. cGANs are a type of GAN that condition the *Generator_{Network}* and *Discriminator_{Network}* on additional information, s.a. class labels or attribute vectors.

In addition to image synthesis, GANs have been applied to a wide range of problems, s.a. data augmentation, style transfer [9], and anomaly detection [10]. GAN-based image synthesis has seen important advances in recent years, with the introduction of progressive GANs, styleGAN [11], and BigGAN. These models are able to create high-quality images with high resolution and diverse styles.

III. GAN ARCHITECTURE

GANs [6] are a type of generative model that learn to generate new data samples that resemble a given *TrainingData*. The basic GAN architecture consists of two NNs: a *Generator_{Network}* and a *Discriminator_{Network}*. The *Generator_{Network}* considers a random noise vector $v \in \mathbb{R}^d$ as input and creates a synthetic data (*SyntheticData*)




Fig. 2. The general structure of a GAN.

sample $\hat{s} \in \mathbb{R}^m$ as output. The *DiscriminatorNetwork* takes a data sample $s \in \mathbb{R}^m$ as input and produces a scalar value $D(s) \in [0, 1]$ as output, indicating the probability that s is a *RealData* sample (as opposed to a synthetic sample produced by the *GeneratorNetwork*).

The *GeneratorNetwork* and The *DiscriminatorNetwork* are trained in an adversarial way, with the *GeneratorNetwork* attempting to generate synthetic samples that resemble *RealData* samples and the *DiscriminatorNetwork* attempting to differentiate between *RealData* and *SyntheticData* samples. The *TrainingProcess* can be modeled as a 2-player minimax game in which the *GeneratorNetwork* G attempts to minimize the following objective function:

$$\min_G \max_D \mathbb{E}_{s \sim p_{\text{data}}(s)} [\log D(s)] + \mathbb{E}_{v \sim p_v(v)} [\log(1 - D(G(v)))], \quad (1)$$

where:

- $p_{\text{data}}(s)$ = true data distribution;
- v = noise vector;
- $p_v(v)$ = prior distribution of v ;
- \mathbb{E} = the expected value.

The first term in Equation 1 encourages the *DiscriminatorNetwork* to correctly classify *RealData* samples as real, while the second term encourages the *GeneratorNetwork* to generate synthetic samples that the *DiscriminatorNetwork* classifies as real.

The *DiscriminatorNetwork* D tries to maximize the following objective function:

$$\max_D \mathbb{E}_{s \sim p_{\text{data}}(s)} [\log D(s)] + \mathbb{E}_{v \sim p_v(v)} [\log(1 - D(G(v)))], \quad (2)$$

where the first term encourages the *DiscriminatorNetwork* to correctly classify *RealData* samples as real, while the second term encourages the *DiscriminatorNetwork* to correctly classify synthetic samples as fake.

Overall, GANs have shown remarkable success in generating high-quality *SyntheticData* samples in a variety of domains, including images, audio, and text. The different types of GAN architectures continue to evolve and improve, and hold great promise for future advancements in generative modeling.

IV. LOSS FUNCTIONS FOR GANS

The success of GANs in generating high-quality *SyntheticData* samples is closely tied to the design of their loss functions (*LossFunction*). In this section, we review some of the most commonly utilized LFs for GANs and their properties.

A. The Original GAN *LossFunction*

The original GAN *LossFunction* [6] is given by Equation 1, which encourages the *GeneratorNetwork* G to generate synthetic samples that are indistinguishable from real samples by the *DiscriminatorNetwork* D . While the original GAN *LossFunction* has been successful in generating high-quality *SyntheticData* samples, it suffers from several problems, including instability during training and mode collapse, where the *GeneratorNetwork* learns to produce a limited set of samples that do not represent the full diversity of the true data distribution.

B. Improved GAN LFs

To address the problems with the original GAN LF, several improved GAN LFs have been proposed in the literature.

Wasserstein GANs (WGANs) [7] use the Wasserstein distance as a LF, which has been shown to produce more stable training and produce higher-quality samples. The WGAN *LossFunction* is given by Equation 3:

$$\min_G \max_D \mathbb{E}_{s \sim p_{\text{data}}(s)} [D(s)] - \mathbb{E}_{v \sim p_v(v)} [D(G(v))], \quad (3)$$

where D is a 1-Lipschitz function, and the *DiscriminatorNetwork* D is trained to maximize Equation 3.

Another approach is to use a least-squares LF, as proposed in the Least Squares GAN (LSGAN) [12]. The LSGAN *LossFunction* is given by Equation 4:

$$\min_G \max_D \frac{1}{2} \mathbb{E}_{s \sim p_{\text{data}}(s)} [(D(s) - 1)^2] + \frac{1}{2} \mathbb{E}_{v \sim p_v(v)} [D(G(v))^2], \quad (4)$$

where the *DiscriminatorNetwork* D is trained to minimize Equation 4.

Other approaches include the hinge *LossFunction* utilized in HingeGANs [13] and the feature matching *LossFunction* utilized in Feature Matching GANs [14].

In conclusion, the choice of $Loss_{Function}$ is critical for the success of GANs in generating high-quality $Synthetic_{Data}$ samples. While the original GAN $Loss_{Function}$ has been successful in many applications, several improved LFs have been proposed that address its limitations and produce more stable training and higher-quality samples. The type of $Loss_{Function}$ used is determined by the purpose and problem at hand.

V. TRAINING GANs

GANs are typically trained using a 2-player minimax game, where a $Generator_{Network}$ learns to produce $Synthetic_{Data}$ samples, and a $Discriminator_{Network}$ learns to differentiate between $Real_{Data}$ and $Synthetic_{Data}$ samples. The $Training_{Process}$ involves iteratively updating the parameters of the $Generator_{Network}$ and $Discriminator_{Network}$ to improve their performance.

A. Challenges in Training GANs

Training GANs can be challenging due to several factors, including instability and mode collapse. Instability can arise when the $Discriminator_{Network}$ is too powerful and quickly learns to differentiate between $Real_{Data}$ and $Synthetic_{Data}$ samples, making it difficult for the $Generator_{Network}$ to learn. Mode collapse, on the other hand, can occur when the $Generator_{Network}$ learns to create a restricted number of samples that fail to accurately reflect the diversity of the underlying data distribution. To address instability, several approaches have been proposed, s.a. modifying the GAN $Loss_{Function}$ to make it more stable during training. For example, the Wasserstein GAN (WGAN) [7] replaces the original GAN $Loss_{Function}$ with the Wasserstein distance, which can produce more stable training. Several strategies have been developed to address mode collapse, s.a. adding noise to the input of the $Generator_{Network}$, using feature matching [14], or using different architectures for the $Generator_{Network}$ and $Discriminator_{Network}$, s.a. the CycleGAN [15].

B. Stabilizing GAN Training

Several strategies for stabilizing GAN training and addressing the aforementioned issues have been proposed. One such technique is minibatch discrimination [14], which involves adding additional features to the $Discriminator_{Network}$ that allow it to compare multiple samples at once and differentiate between them. This enhances the diversity of the produced samples and helps to prevent mode collapse. Another technique is spectral normalization [13], which involves normalizing the weights of the $Discriminator_{Network}$ to ensure that the Lipschitz constant of the network is limited. This helps to prevent the $Discriminator_{Network}$ from becoming too powerful and stabilizes the $Training_{Process}$. Other techniques include using different LFs, s.a. the least-squares GAN (LSGAN) $Loss_{Function}$ [12], to improve the stability of the $Training_{Process}$. Additionally, regularization techniques, s.a. weight decay and dropout, can also be utilized to prevent overfitting and improve the generalization performance of the models.

Consequently, training GANs is a challenging task which needs careful consideration of several factors to achieve stable and high-quality results. The key challenges in GAN training include instability and mode collapse, which can be addressed using various techniques, s.a. modifying the GAN LF, using different architectures, and adding regularization. Further research is needed to develop more effective techniques for training GANs and improving their performance in various applications.

VI. APPLICATIONS OF GANs

GANs have gotten a lot of interest recently because of their capacity to produce high-quality $Synthetic_{Data}$ that closely matches $Real_{Data}$. GANs have various applications in different fields, including computer vision, NLP, and healthcare.

A. Image Synthesis

One of the most famous applications of GANs is in image synthesis, where GANs are utilized to produce new images that are similar to a given set of training images. GANs can create highly realistic images that can be utilized for various purposes, s.a. in video games, virtual reality, and creating $Synthetic_{Data}$ for training ML models. Recent advances in GAN-based image synthesis have led to the development of several new techniques s.a. progressive GANs, styleGAN [16], and BigGAN [17]. Progressive GANs produce high-resolution images by incrementally increasing the size of the produced images, while styleGAN allows for the control of different aspects of the produced images s.a. style, pose, and facial expression. BigGAN is capable of generating high-quality images with up to 512x512 pixels.

B. Data Augmentation

GANs are additionally appropriate for data augmentation, which involves creating $Synthetic_{Data}$ to expand the size of the $Training_{Data}$. By supplying additional $Training_{Data}$ that is close to the $Real_{Data}$, data augmentation using GANs may increase the performance of ML models. This approach has been successfully applied in various areas s.a. object detection, image classification, and speech recognition [18].

C. Style Transfer

Additionally, GANs can be used for style transfer, which is the process of transferring an image's style to another. This method can be used to create novel pictures by merging the content of one picture with the style of another. Style transfer using GANs has shown promising results in various domains, s.a. fashion design, art, and photography [19].

D. Emerging Applications

GANs are also being utilized for emerging applications s.a. video synthesis and text-to-image synthesis. In video synthesis, GANs are utilized to produce new video frames that are similar to the existing frames. This technique can be utilized to create high-quality videos with less manual effort. Text-to-image synthesis using GANs involves generating images from textual descriptions. This approach has potential applications

in fashion design, interior design, and other areas where the ability to produce images from textual descriptions can be useful [20].

VII. EVALUATION OF GANS

Due to the lack of a clear goal function, evaluating the performance of GANs is a difficult undertaking. GANs produce *SyntheticData* by learning the underlying distribution of the *TrainingData*, and the quality of the produced data depends on different factors s.a. the architecture of the *GeneratorNetwork* and *DiscriminatorNetwork*, the optimization algorithm utilized, and the choice of hyperparameters. To evaluate the performance of GANs, various metrics have been defined, including IS⁴ and FID⁵. Based on the classification precision of a previously trained Inception model, the IS assesses the diversity and quality of the generated images. Using an Inception model that has already been trained, the FID calculates the separation in feature space between the distributions of the real and created images.

While these metrics have been widely utilized in GAN research, they have limitations. For example, the IS is known to favor models that produce images that are easily classified by the Inception model, even if they are low-quality or lack diversity. The FID can be sensitive to noise and image artifacts, and may not always correlate with visual quality. Moreover, both metrics require pre-trained Inception models, which may not be readily available or may not be suitable for all types of data. There is ongoing research to develop new evaluation metrics that can better capture the performance of GANs. Some recent proposals include KID⁶, which measures the distance between the distributions of the features extracted from the Inception model, and LPIPS⁷, which measures the perceptual similarity between images based on the activations of a pre-trained deep NN.

In conclusion, evaluating the performance of GANs is an important aspect of GAN research, and various metrics have been proposed for this purpose. However, current evaluation metrics have limitations, and there is a need for new metrics that can better capture the performance of GANs.

VIII. CHALLENGES AND OPEN ISSUES

GANs have demonstrated enormous potential for producing realistic photos, movies, and other forms of data. However, there are various obstacles and unresolved concerns that must be addressed in order to increase the performance and usability of GANs. Among the obstacles and unresolved issues are:

- 1) Mode collapse: Mode collapse is a common problem in GANs, where the *GeneratorNetwork* creates only a limited set of outputs, ignoring other possible outputs. This might lead to a lack of diversity in the data that is generated. One possible cause of mode collapse is the *DiscriminatorNetwork* being too strong

compared to the *GeneratorNetwork*, which leads to the *GeneratorNetwork* outputting similar samples that fool the *DiscriminatorNetwork*. Researchers are exploring various techniques to address mode collapse, s.a. adding regularization terms to the *LossFunction* or using alternative training methods. Other techniques include modifying the architecture of the *GeneratorNetwork* and *DiscriminatorNetwork* or using more advanced optimization methods.

- 2) Training instability: It can be challenging to train GANs, and the *TrainingProcess* can be unstable, leading to oscillations or divergence in the *GeneratorNetwork* and *DiscriminatorNetwork* loss. This can make it challenging to achieve good performance. One possible cause of training instability is the imbalance between the *GeneratorNetwork* and *DiscriminatorNetwork*, where one dominates the other. Another cause is the vanishing gradient problem, where the gradients of the *LossFunction* become too small to update the parameters. Researchers are investigating various approaches to improve the stability of GAN training, s.a. adjusting the learning rate, using different optimization algorithms, or adding noise to the *TrainingProcess*. Another approach is to use more advanced architectures, s.a. Wasserstein GANs or Spectral Normalization GANs, that have been shown to be more stable during training.
- 3) Evaluation metrics: There is a lack of widely accepted evaluation metrics for GANs, making it difficult to compare different models and assess their performance objectively. Some proposed evaluation metrics include FID, which evaluates the distance between the distribution of produced samples and the distribution of real samples in a feature space, and IS, which measures the diversity and quality of produced samples based on their classification scores by a pre-trained classifier. However, these metrics have limitations and may not capture all aspects of the produced data. Researchers are exploring alternative evaluation metrics and methods to better quantify the performance of GANs.
- 4) Scalability: GANs can be computationally expensive to train and require enormous volumes of data. As a result, its scalability and applicability to real-world challenges may be limited. Transfer learning, which involves fine-tuning a pre-trained model on a new dataset, is one method for improving GAN scalability. Another option is to leverage parallel computing, also known as distributed training over numerous GPUs, or to use cloud-based computing resources. Researchers are also looking into techniques to reduce the quantity of data necessary for GAN training, also known as semi-supervised learning, or to enrich the *TrainingData* with generative models.
- 5) Ethical implications: As with any technology, GANs raise ethical implications, particularly in the context of generating realistic images or videos. GANs can be utilized to create fake content that can be utilized for ma-

⁴IS=Inception Score

⁵FID= Fréchet Inception Distance

⁶KID=Kernel Inception Distance

⁷LPIPS=Learned Perceptual Image Patch Similarity

TABLE I
CHALLENGES AND OPEN ISSUES IN GANS

Challenge/Open Issue	Description
Mode Collapse	Mode collapse, in which the GTR only generates a small number of outputs, might cause GANs to lose diversity in the data they generate. To address mode collapse, researchers are investigating numerous approaches.
Training Instability	The training of GANs can be challenging and unstable, which can cause oscillations or divergence in the GTR and DTR loss. To increase the stability of GAN training, researchers are looking into a number of different strategies.
Evaluation Metrics	There is a lack of widely accepted evaluation metrics for GANs, making it difficult to compare different models and assess their performance objectively. Researchers are exploring alternative evaluation metrics and methods.
Scalability	GANs' scalability and suitability for use in solving real-world problems are constrained by their computationally expensive and data-intensive training requirements. Researchers are looking into techniques to reduce the quantity of data necessary for GAN training or to leverage parallel computing.
Ethical Implications	GANs raise ethical implications, particularly in the context of generating realistic images or videos. Researchers and policymakers are exploring ways to mitigate the risks and promote responsible use of GANs.

licious purposes, s.a. spreading disinformation or generating deepfakes. This can have serious consequences for individuals and society as a whole. Researchers and policymakers are exploring ways to mitigate these risks and promote responsible use of GANs, s.a. developing detection methods for deepfakes or creating guidelines for the ethical use of *SyntheticData*.

Table I summarizes the challenges and open issues in GANs. Overall, addressing these challenges and open issues will be critical for realizing the full potential of GANs and ensuring their responsible and ethical use.

IX. FUTURE DIRECTIONS

GANs have rapidly become one of the most exciting fields in deep learning since their introduction in 2014. The ability to generate realistic data using GANs has numerous applications in various domains such as computer vision, natural language processing, and audio synthesis. As GANs continue to evolve, researchers and practitioners can explore several future directions to improve their scalability, stability, and performance. These future directions include:

- 1) Improving the scalability of GANs: GANs can be computationally expensive and require large amounts of data to train. Researchers are exploring various techniques to improve the scalability of GANs, s.a. using parallel computing, transfer learning, or reducing the amount of data required for GAN training. Practitioners can experiment with these techniques to improve the scalability of their GAN models and make them more applicable to real-world problems.

- 2) Developing more advanced architectures: GAN architectures have evolved significantly since their inception, from the original GAN architecture to more advanced architectures s.a. Wasserstein GANs, Progressive GANs, and StyleGANs. Researchers can continue to explore and develop new architectures that are more stable, scalable, and capable of generating high-quality data.
- 3) Incorporating domain knowledge: GANs can benefit from incorporating domain knowledge, s.a. physical laws or expert knowledge in a particular field. Researchers can explore ways to incorporate domain knowledge into GAN models to improve their performance and generalization capabilities.
- 4) Exploring new applications of GANs: GANs have already been applied to a wide range of domains, s.a. computer vision, NLP, and audio synthesis. However, there are still many potential applications of GANs that have yet to be explored, s.a. in healthcare, finance, or social sciences. Researchers can explore new use cases and applications of GANs in these domains, and practitioners can experiment with applying GANs to new problems in their field.
- 5) Developing more robust evaluation metrics: There is a need for more robust evaluation metrics for GANs that can capture all aspects of the produced data, s.a. visual quality, diversity, and realism. Researchers can explore new evaluation metrics and methods that are more reliable and objective, and practitioners can use these metrics to estimate the performance of their GAN models.
- 6) Use of formal methods: Another future direction for GANs is the use of formal methods to improve their reliability and safety [21], [22]. Formal methods are a set of mathematical techniques and tools used to rigorously analyze and verify software and hardware systems. The use of formal methods in GANs can help ensure that they produce outputs that meet certain safety and reliability requirements. Formal methods can be used to verify properties of GANs, such as the absence of certain types of errors or the correctness of certain operations. For example, formal methods can be used to verify that the produced data does not violate certain safety constraints, or that the *TrainingProcess* does not diverge or exhibit undesirable behavior.

As summarized in Table II, there are many exciting future directions for GANs, including improving the scalability and stability of GANs, developing more advanced architectures, incorporating domain knowledge, exploring new applications of GANs, and developing more robust evaluation metrics. By continuing to push the boundaries of GAN research and development, researchers and practitioners can unlock the full potential of GANs and drive innovation in a wide range of domains.

TABLE II
FUTURE DIRECTIONS OF GANS

Direction	Description
Improving Scalability	Researchers are exploring various techniques to improve the scalability of GANs, s.a. using parallel computing, transfer learning, or reducing the amount of data required for GAN training.
Developing Advanced Architectures	Researchers can continue to explore and develop new architectures that are more stable, scalable, and capable of generating high-quality data.
Incorporating Domain Knowledge	GANs can benefit from incorporating domain knowledge, s.a. physical laws or expert knowledge in a particular field. Researchers can explore ways to incorporate domain knowledge into GAN models.
Exploring New Applications	Researchers can explore new use cases and applications of GANs in domains such as healthcare, finance, or social sciences. Practitioners can experiment with applying GANs to new problems in their field.
Developing Robust Evaluation Metrics	Researchers can explore new evaluation metrics and methods that are more reliable and objective, and practitioners can use these metrics to estimate the performance of their GAN models.
Use of Formal Methods	Formal methods can be used to improve the reliability and safety of GANs by verifying the properties of GANs or generating adversarial examples to evaluate their robustness [23], [24].

X. CONCLUSION

GANs have emerged as a potent tool for producing realistic data in a variety of disciplines. In this paper, we have proposed a comprehensive guide to GANs, covering their architecture, LFs, training methods, applications, evaluation metrics, challenges, and future directions. We have reviewed the historical development of GANs, from their original formulation to more recent advances s.a. Wasserstein GANs and StyleGANs. We have discussed the key design choices and variations in the GAN architecture, as well as the different LFs utilized to train GAN models. We have also explored the various applications of GANs, from image synthesis to NLP and audio synthesis, and reviewed the evaluation metrics utilized to assess the diversity and quality of GAN-produced data. Additionally, we have highlighted the challenges and open issues in GAN research, s.a. training instability, mode collapse, and ethical considerations. Finally, we have provided a glimpse into the future directions of GAN research, including improving scalability, developing new architectures, incorporating domain knowledge, and exploring new applications.

Overall, GANs represent a rapidly evolving field of research with tremendous potential for innovation and impact. By providing a comprehensive guide to GANs, we hope to facilitate further research and development in this area, and inspire new applications and use cases. We anticipate that as GANs improve, they will play a growing importance in data production and synthesis, opening up new avenues for scientific discovery and artistic expression.

REFERENCES

- [1] M. Krichen, A. Mihoub, M. Y. Alzahrani, W. Y. H. Adoni, and T. Nahhal, "Are formal methods applicable to machine learning and artificial intelligence?" in *2022 2nd International Conference of Smart Systems and Emerging Technologies (SMARTTECH)*. IEEE, 2022, pp. 48–53.
- [2] W. Boulila, M. Driss, E. Alshanqiti, M. Al-Sarem, F. Saeed, and M. Krichen, "Weight initialization techniques for deep learning algorithms in remote sensing: Recent trends and future perspectives," *Advances on Smart and Soft Computing: Proceedings of ICACIn 2021*, pp. 477–484, 2022.
- [3] E. Brophy, Z. Wang, Q. She, and T. Ward, "Generative adversarial networks in time series: A systematic literature review," *ACM Computing Surveys*, vol. 55, no. 10, pp. 1–31, 2023.
- [4] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *nature*, vol. 521, no. 7553, pp. 436–444, 2015.
- [5] I. Goodfellow, Y. Bengio, and A. Courville, *Deep learning*. MIT press, 2016.
- [6] I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative adversarial networks," *Communications of the ACM*, vol. 63, no. 11, pp. 139–144, 2020.
- [7] M. Arjovsky, S. Chintala, and L. Bottou, "Wasserstein generative adversarial networks," in *International conference on machine learning*. PMLR, 2017, pp. 214–223.
- [8] J. Gauthier, "Conditional generative adversarial nets for convolutional face generation," *Class project for Stanford CS231N: convolutional neural networks for visual recognition, Winter semester*, vol. 2014, no. 5, p. 2, 2014.
- [9] L. A. Gatys, A. S. Ecker, and M. Bethge, "Image style transfer using convolutional neural networks," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2016, pp. 2414–2423.
- [10] X. Xia, X. Pan, N. Li, X. He, L. Ma, X. Zhang, and N. Ding, "Gan-based anomaly detection: a review," *Neurocomputing*, 2022.
- [11] T. Karras, S. Laine, and T. Aila, "A style-based generator architecture for generative adversarial networks," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2019, pp. 4401–4410.
- [12] X. Mao, Q. Li, H. Xie, R. Y. Lau, Z. Wang, and S. Paul Smolley, "Least squares generative adversarial networks," in *Proceedings of the IEEE international conference on computer vision*, 2017, pp. 2794–2802.
- [13] C. Xiaopeng, C. Jiangzhong, L. Yuqin, and D. Qingyun, "Improved training of spectral normalization generative adversarial networks," in *2020 2nd World Symposium on Artificial Intelligence (WSAI)*. IEEE, 2020, pp. 24–28.
- [14] T. Salimans, I. Goodfellow, W. Zaremba, V. Cheung, A. Radford, and X. Chen, "Improved techniques for training gans," *Advances in neural information processing systems*, vol. 29, 2016.
- [15] J.-Y. Zhu, T. Park, P. Isola, and A. A. Efros, "Unpaired image-to-image translation using cycle-consistent adversarial networks," in *Proceedings of the IEEE international conference on computer vision*, 2017, pp. 2223–2232.
- [16] T. Karras, S. Laine, and T. Aila, "A style-based generator architecture for generative adversarial networks," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2019, pp. 4401–4410.
- [17] A. Brock, J. Donahue, and K. Simonyan, "Large scale gan training for high fidelity natural image synthesis," *arXiv preprint arXiv:1809.11096*, 2018.
- [18] C. Shorten and T. M. Khoshgoftaar, "A survey on image data augmentation for deep learning," *Journal of big data*, vol. 6, no. 1, pp. 1–48, 2019.
- [19] L. A. Gatys, A. S. Ecker, and M. Bethge, "Image style transfer using convolutional neural networks," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 2414–2423.
- [20] S. Reed, Z. Akata, X. Yan, L. Logeswaran, B. Schiele, and H. Lee, "Generative adversarial text to image synthesis," in *International conference on machine learning*. PMLR, 2016, pp. 1060–1069.
- [21] M. Krichen, "Contributions to model-based testing of dynamic and distributed real-time systems," Ph.D. dissertation, École Nationale d'Ingénieurs de Sfax (Tunisie), 2018.
- [22] M. Krichen and S. Tripakis, "Interesting properties of the real-time conformance relation tioco," in *Theoretical Aspects of Computing-ICTAC 2006: Third International Colloquium, Tunis, Tunisia, November 20-24, 2006. Proceedings 3*. Springer Berlin Heidelberg, 2006, pp. 317–331.
- [23] M. Krichen, "A formal framework for conformance testing of distributed real-time systems," in *International Conference On Principles Of Distributed Systems*. Springer, 2010, pp. 139–142.
- [24] M. Krichen and S. Tripakis, "State identification problems for timed automata," in *TestCom*, vol. 5, 2005, pp. 175–191.